


A second camera helps in proportion to how differently it looks at the robot


Real captures throughout: the balanced set-pose grid, 1844 clear detections, 575 camera pairs that see the same held pose, current floor-plane projection. The error shapes are pure geometry — one pixel-noise number, nothing fitted. Ground truth is the commanded pose and only ever scores.

A camera does not measure x and y equally well
each error ellipse is stretched along that camera's own sightline




x (m) · ellipses are 1 px of detector error, drawn 22x oversized
geometry says 1.7–1.8x longer than wide; the real residuals measure 1.5–1.8x

Two laws, disagreeing about where the sweet spot is
the random part dips in the middle; the lean falls all the way to zero



The second camera cancels the lean, not the noise
and it can only cancel it if it looks from the other side



What a second camera actually delivers
nothing at all until the lean is dealt with first

