Camera A — deployed-corrected residual

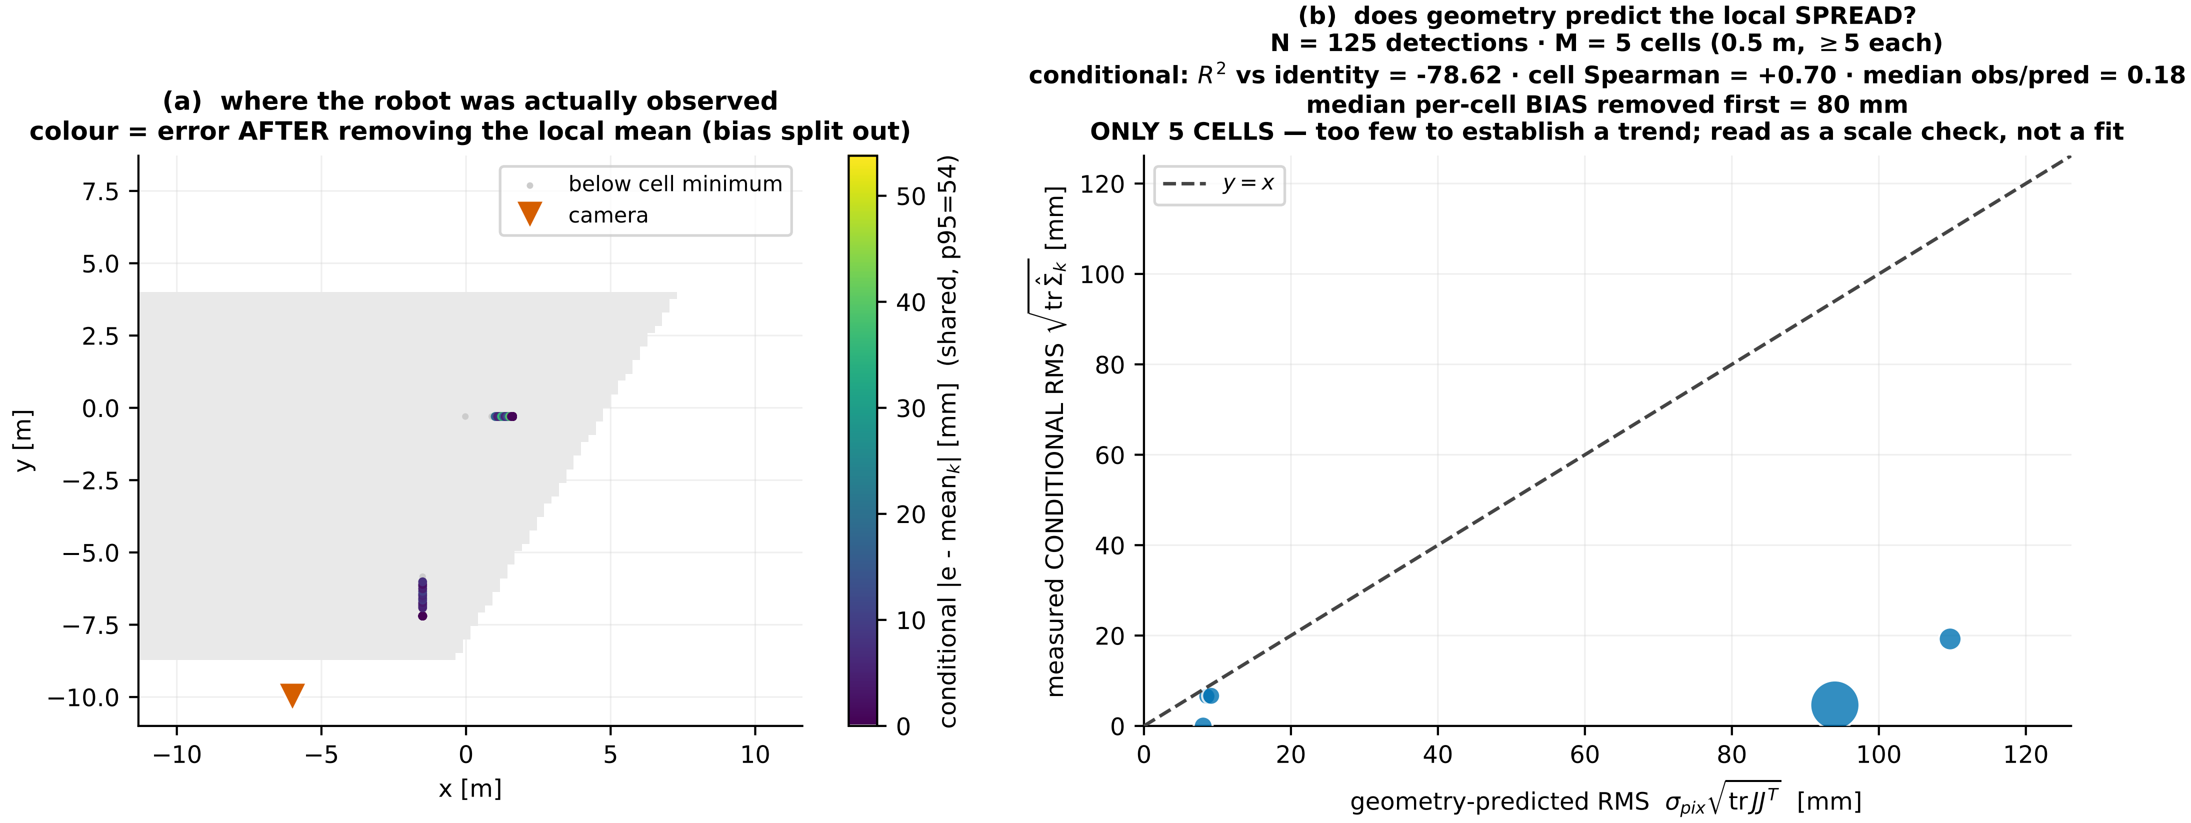

Point area in (b) scales with detections in the cell. sigma_pix is fitted on leave-region-out folds, so each cell's prediction is out-of-region. Ground truth positions the dots in (a) and measures the residual; it never enters a projection, Jacobian or covariance. READ THE LIMIT BEFORE THE RESULT: only 3.1 x 6.9 m is sampled and image column tracks range at |rho| = 0.91, so range, image row and projection conditioning are collinear here. A weak or negative R^2 does NOT show that geometry fails -- it shows this dataset cannot test it.