



Camera A — deployed-corrected residual



(b) does geometry predict the local error?
 N = 125 detections · M = 5 cells (0.5 m, ≥ 5 each)
 held-out R^2 vs identity = -5.72 · cell Spearman = -0.80 · median obs/pred = 8.79
ONLY 5 CELLS — too few to establish a trend; read as a scale check, not a fit



Point area in (b) scales with detections in the cell. sigma_pix is fitted on leave-region-out folds, so each cell's prediction is out-of-region. Ground truth positions the dots in (a) and measures the residual; it never enters a projection, Jacobian or covariance. The sampled ribbon in (a) is the binding limit of this dataset: the footprint is not covered in two dimensions, so range and the full Jacobian cannot be separated here.