



Camera B — deployed-corrected residual



(b) does geometry predict the local error?
N = 295 detections · M = 7 cells (0.5 m, ≥ 5 each)
held-out R^2 vs identity = -1.50 · cell Spearman = +0.18 · median obs/pred = 0.95
ONLY 7 CELLS — too few to establish a trend; read as a scale check, not a fit



Point area in (b) scales with detections in the cell. σ_{pix} is fitted on leave-region-out folds, so each cell's prediction is out-of-region. Ground truth positions the dots in (a) and measures the residual; it never enters a projection, Jacobian or covariance. The sampled ribbon in (a) is the binding limit of this dataset: the footprint is not covered in two dimensions, so range and the full Jacobian cannot be separated here.