Camera B — deployed path: bottom pixel @ $z = 0.05$ m + v2 along-bearing fit

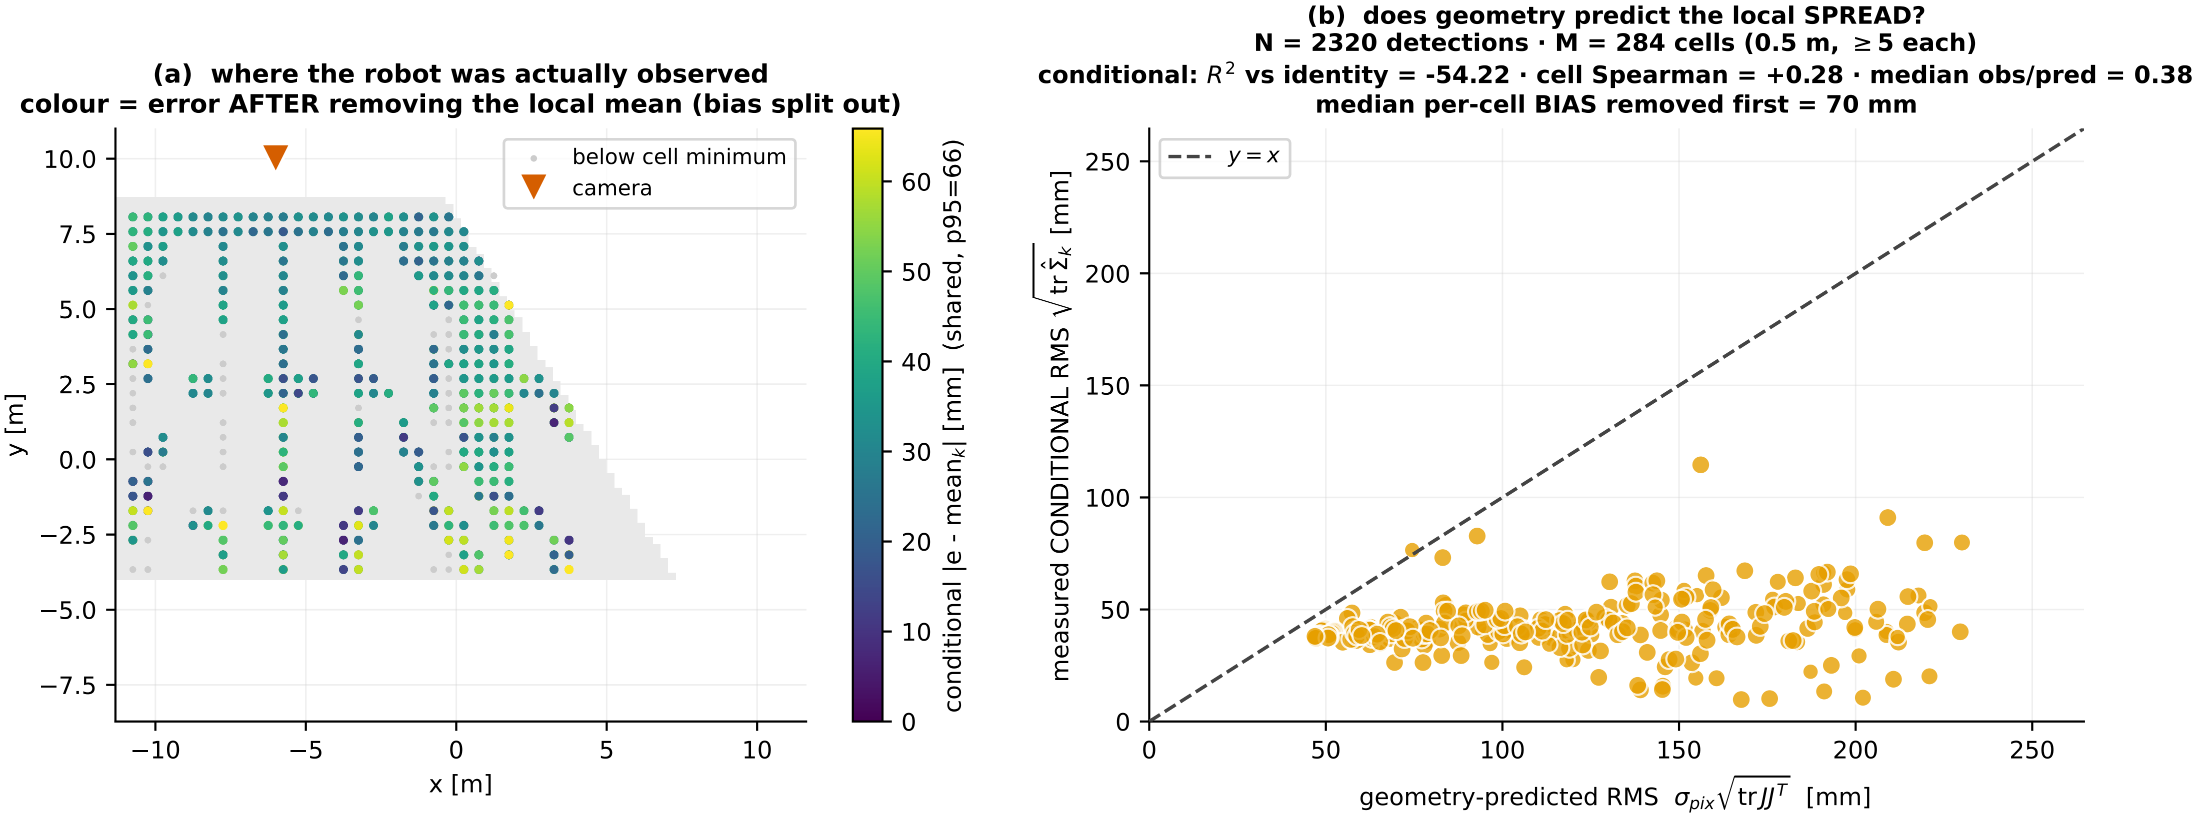

Point area in (b) scales with detections in the cell. sigma_pix is fitted on leave-region-out folds, so each cell's prediction is out-of-region. Ground truth positions the dots in (a) and measures the residual; it never enters a projection, Jacobian or covariance. COVERAGE: 14.5 x 11.7 m sampled in two dimensions across 8 headings; image column is decorrelated from range (|rho| = 0.27), so this dataset CAN separate image position from distance. It still cannot separate sigma_max(J) from range: those stay collinear at |rho| ~ 0.98 because for an oblique ground-plane camera the amplification is set by depth, which is what range measures. That is a property of the geometry, not of the sampling.