**EXP-RCOND — NULL: a per-camera conditional covariance does not beat one constant pooled matrix**
**$n$ = 1425 associated detections, 4 cameras, 3 captures · no ground truth enters either arm**

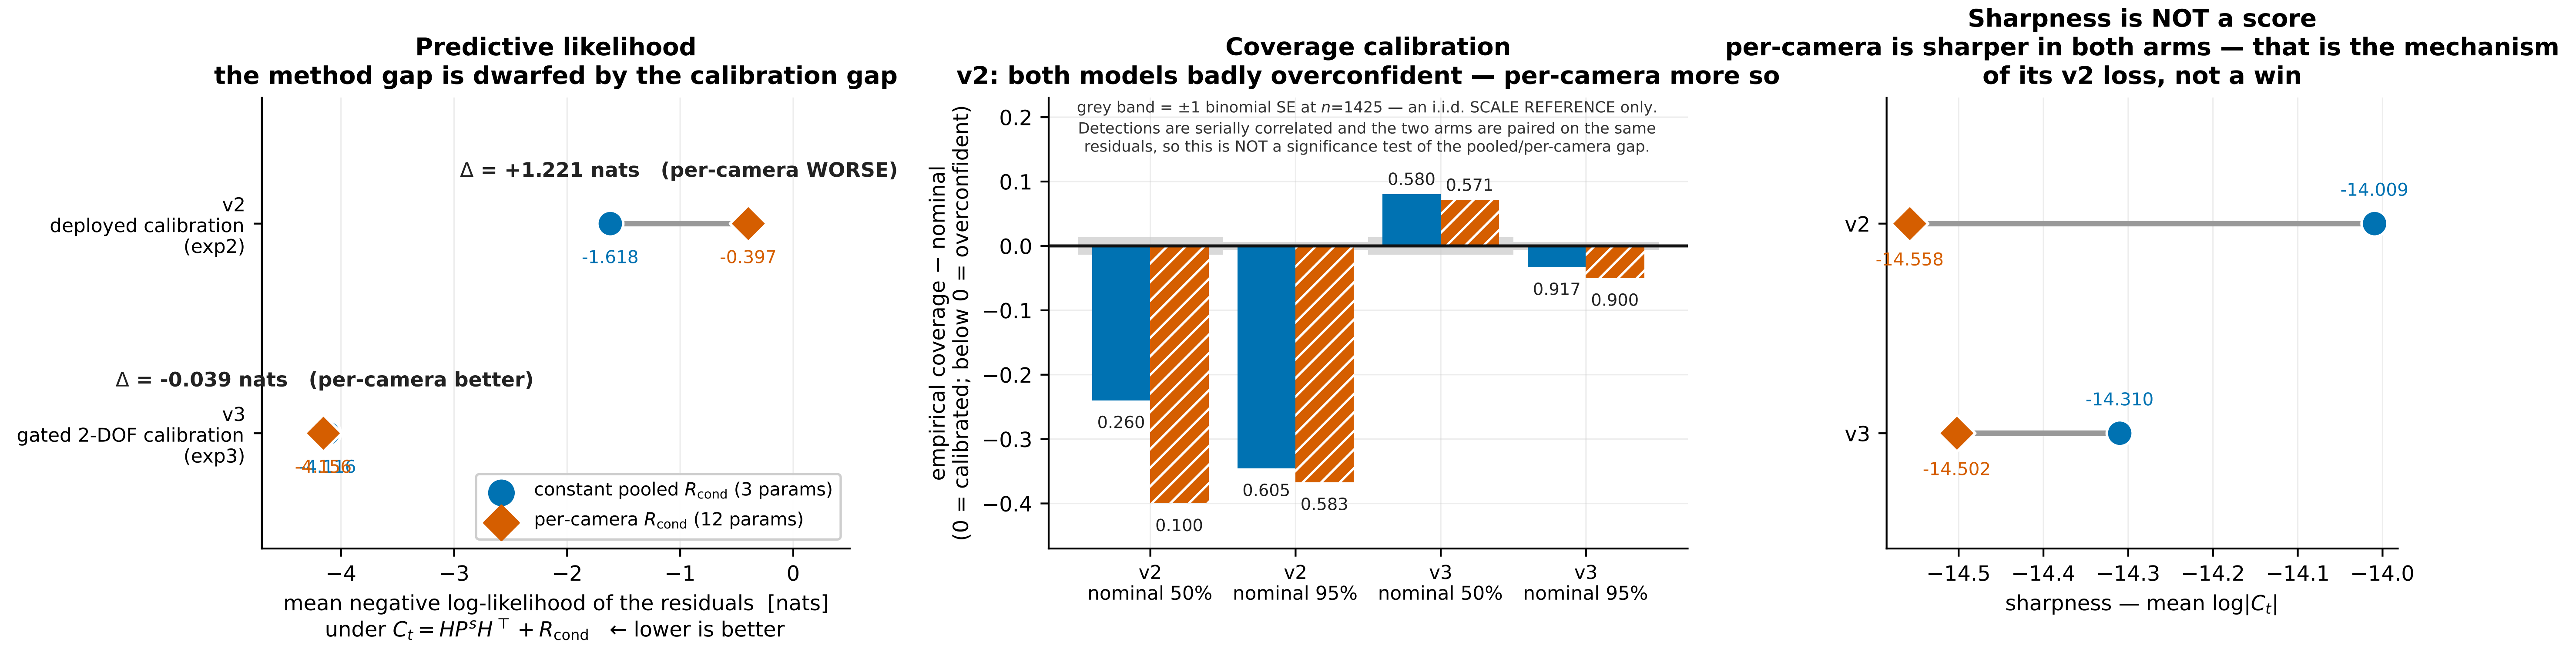


Read with the caveat, not around it: both arms are scored on the SAME 1425 residuals that fitted them. “Held-out” in this study names the leave-one-camera-out reference trajectory, not held-out scoring data.
The comparison is therefore biased TOWARD per-camera — 12 free parameters against 3, scored in-sample — and per-camera still loses by 1.221 nats under v2 and gains 0.039 nats under v3 while losing 95% coverage. That is a tie, not a promotion.